

**fuse** | machines



BONE  
FRACTURE  
DETECTION

FUSE AI FELLOWSHIP  
2023-2024



# 1. TEAM MEMBERS

a) Joaquin Ayzanoa (Perú)



## 2. PROBLEM STATEMENT

- a) A data set of bone fractures with which we want to train a machine learning model, specifically focusing on algorithms to automatically detect and classify bone fractures in X-ray images for better and fast registration for doctors and patients.



# 3. OBJECTIVES

- a) Detect and classify bone fractures in X-ray images for better and fast registration for doctors and patients.
- b) Deploy the ML model with an interface so it can predict the type of fracture when someone upload an image.



# 4. GITHUB REPOSITORY LINK

This is the link of the repository:

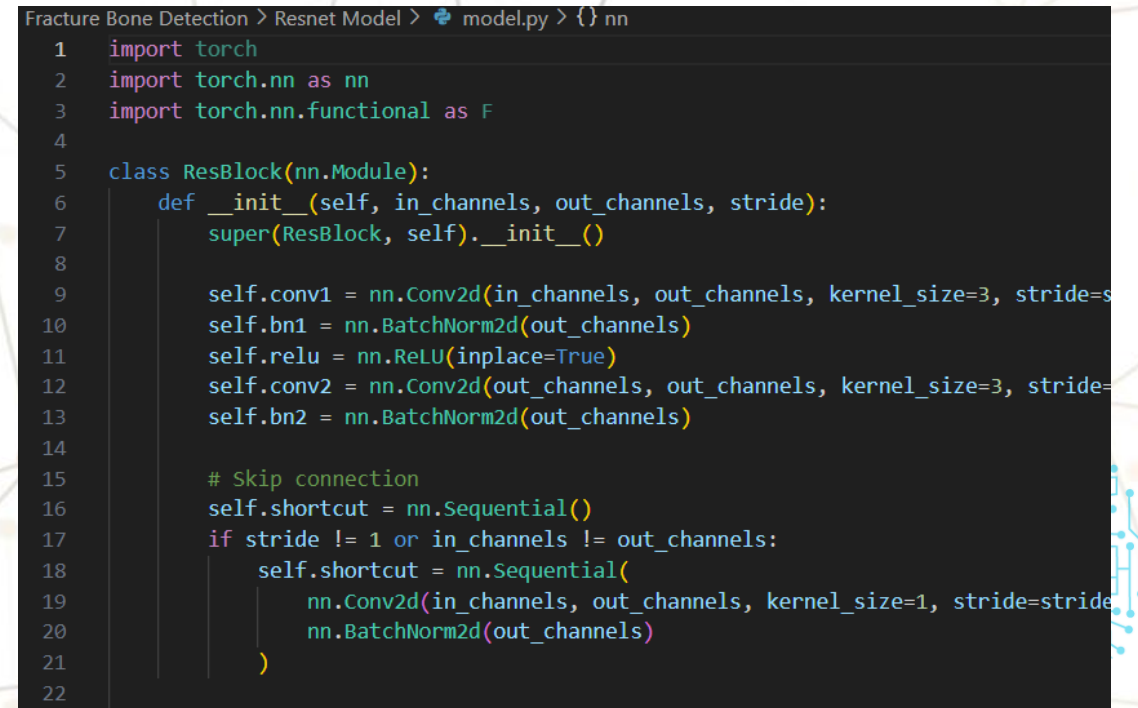
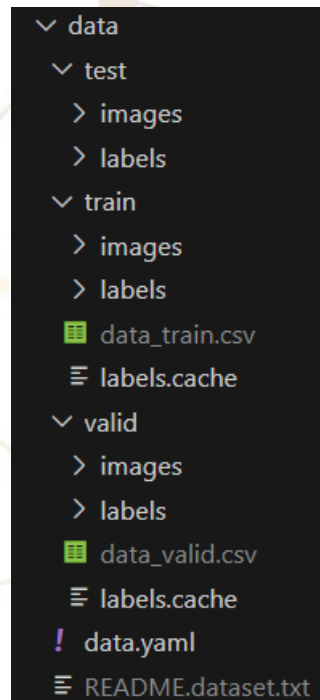
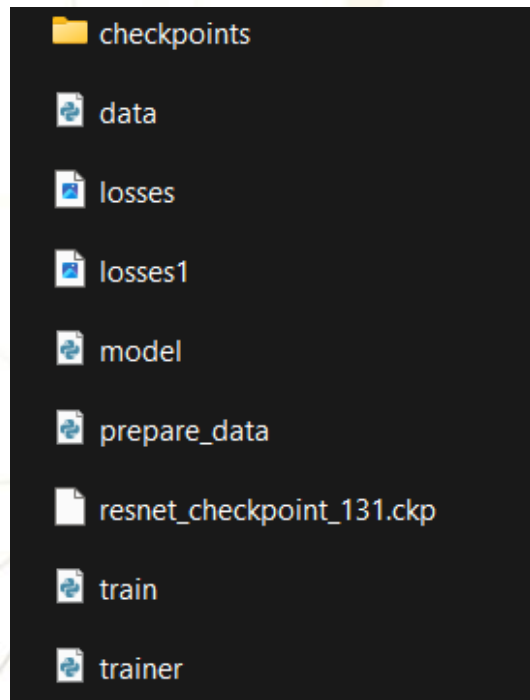
<https://github.com/JoaquinAyzanoa/PortafolioJA/tree/e96351fa14b03b0788fd2a7d78b996a5b3a9c38d/Fracture%20Bone%20Detection>





# 5. RESNET MODEL

The Resnet model was created from scratch using pytorch, no pretrained model or downloaded model was used here.




# 5. RESNET MODEL

The Resnet model was created from scratch using pytorch, no pretrained model or downloaded model was used here.

```
Fracture Bone Detection > Resnet Model > prepare_data.py > {} pd
1 import glob
2 import os
3 import pandas as pd
4
5 def prepare_data(root, mode):
6     file_dataname = os.path.join(root, mode)
7     file_dataname = file_dataname + f'/data_{mode}.csv'
8     if os.path.exists(file_dataname):
9         df = pd.read_csv(file_dataname, delimiter=';')
10        print(f'data_{mode}.csv' + ' already exists')
11    else:
12        files = sorted(glob.glob(os.path.join(root, mode) + '/images/*.jpg'))
13        df= pd.DataFrame(columns=['filename','label', 'coords'])
14        for file in files:
15            text_file = file.replace(".jpg", ".txt").replace("images", "labels")
16
17            with open(text_file, mode="r") as f:
18                lines = f.readlines()
19                for line in lines:
20                    values = [float(value) for value in line.split()]
21                    label = int(values[0])
22                    coords = values[1:]
23                    row_data = [file]
24                    row_data.append(label)
```

```
Fracture Bone Detection > Resnet Model > train.py > ...
22
23 # set up data loading for the training and validation set each using t.utils.data
24 train_dataset = ChallengeDataset(train_data, mode='train')
25 val_dataset = ChallengeDataset(val_data, mode='val')
26 train_dataloader = t.utils.data.DataLoader(train_dataset, batch_size=48, shuffle=
27 val_dataloader = t.utils.data.DataLoader(val_dataset, batch_size=48, shuffle=False)
28
29 # create an instance of our ResNet model
30 resnet_model = model.ResNet()
31 resnet_model.to(device)
32 # set up a suitable loss criterion (you can find a pre-implemented loss functions
33 # set up the optimizer (see t.optim)
34 # create an object of type Trainer and set its early stopping criterion
35 criterion = t.nn.CrossEntropyLoss() # Assuming it's a multi class classification
36 optimizer = t.optim.Adam(resnet_model.parameters(), lr=0.00001)
37 trainer = Trainer(model=resnet_model, crit=criterion, optim=optimizer,
38                  train_dl=train_dataloader, val_test_dl=val_dataloader,
39                  cuda=True, early_stopping_patience=25)
40
41 #trainer.restore_checkpoint(79)
42 res = trainer.fit(epochs=200)
43
```

# 5.1 Results



## ✓ TERMINAL

```
Epoch 130/200: Train Loss: 1.2447, Val Loss: 1.3684, F1 Score: 0.7963
Epoch 131/200: Train Loss: 1.2428, Val Loss: 1.3483, F1 Score: 0.8056
Epoch 132/200: Train Loss: 1.2381, Val Loss: 1.3751, F1 Score: 0.7709
Epoch 133/200: Train Loss: 1.2427, Val Loss: 1.3700, F1 Score: 0.7849
Epoch 134/200: Train Loss: 1.2407, Val Loss: 1.3798, F1 Score: 0.7667
Epoch 135/200: Train Loss: 1.2460, Val Loss: 1.4967, F1 Score: 0.6689
Epoch 136/200: Train Loss: 1.2396, Val Loss: 1.3573, F1 Score: 0.7905
Epoch 137/200: Train Loss: 1.2312, Val Loss: 1.3677, F1 Score: 0.7814
Epoch 138/200: Train Loss: 1.2327, Val Loss: 1.3985, F1 Score: 0.7460
Epoch 139/200: Train Loss: 1.2372, Val Loss: 1.3769, F1 Score: 0.7764
Epoch 140/200: Train Loss: 1.2328, Val Loss: 1.3675, F1 Score: 0.7830
Epoch 141/200: Train Loss: 1.2302, Val Loss: 1.4016, F1 Score: 0.7574
Epoch 142/200: Train Loss: 1.2360, Val Loss: 1.3950, F1 Score: 0.7550
Epoch 143/200: Train Loss: 1.2355, Val Loss: 1.3733, F1 Score: 0.7756
Epoch 144/200: Train Loss: 1.2303, Val Loss: 1.4430, F1 Score: 0.7138
Epoch 145/200: Train Loss: 1.2395, Val Loss: 1.3915, F1 Score: 0.7705
Epoch 146/200: Train Loss: 1.2372, Val Loss: 1.3624, F1 Score: 0.7839
Epoch 147/200: Train Loss: 1.2335, Val Loss: 1.3638, F1 Score: 0.7819
Epoch 148/200: Train Loss: 1.2293, Val Loss: 1.4029, F1 Score: 0.7375
Epoch 149/200: Train Loss: 1.2357, Val Loss: 1.4042, F1 Score: 0.7306
Epoch 150/200: Train Loss: 1.2290, Val Loss: 1.3804, F1 Score: 0.7608
Epoch 151/200: Train Loss: 1.2305, Val Loss: 1.4086, F1 Score: 0.7491
Epoch 152/200: Train Loss: 1.2302, Val Loss: 1.3770, F1 Score: 0.7741
Epoch 153/200: Train Loss: 1.2270, Val Loss: 1.3767, F1 Score: 0.7921
Epoch 154/200: Train Loss: 1.2288, Val Loss: 1.3628, F1 Score: 0.7973
Epoch 155/200: Train Loss: 1.2262, Val Loss: 1.3854, F1 Score: 0.7686
Epoch 156/200: Train Loss: 1.2275, Val Loss: 1.3872, F1 Score: 0.7690
Early stopping after 156 epochs without improvement.
```



## 5.1 Results


### Fracture Bone Detection

Choose a file

Drag and drop file here  
Limit 200MB per file • JPG, JPEG, PNG


Browse files

coronoid-process-fracture\_jpg.rf.71650459c69a9734ecd545067cf18bf4.jpg 118.6KB ×



Uploaded file

Prediction: Elbow positive





# 6. YOLO-V8 FOR OBJECT DETECTION

The YoloV8 model is a pretrained model downloaded from the library “Ultralytics”.  
The YoloV8 was finetuned to bone fracture detection task.

```
# Run inference on an image with YOLO
!yolo detect train data='./data.yaml' model=yolov8x.pt epochs=50 imgsz=320 device=0
```

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
49/50	5.51G	1.381	0.975	1.226	11	320: 100% 227/227 [01:2
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% 11
	all	348	204	0.309	0.26	0.237 0.0789

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
50/50	5.5G	1.325	0.9576	1.197	7	320: 100% 227/227 [01:2
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% 11
	all	348	204	0.357	0.25	0.22 0.0749

50 epochs completed in 1.312 hours.  
Optimizer stripped from runs/detect/train/weights/last.pt, 136.7MB  
Optimizer stripped from runs/detect/train/weights/best.pt, 136.7MB

Validating runs/detect/train/weights/best.pt...

Ultralytics YOLOv8.1.27 Python-3.10.12 torch-2.1.0+cu121 CUDA:0 (Tesla T4, 15102MiB)

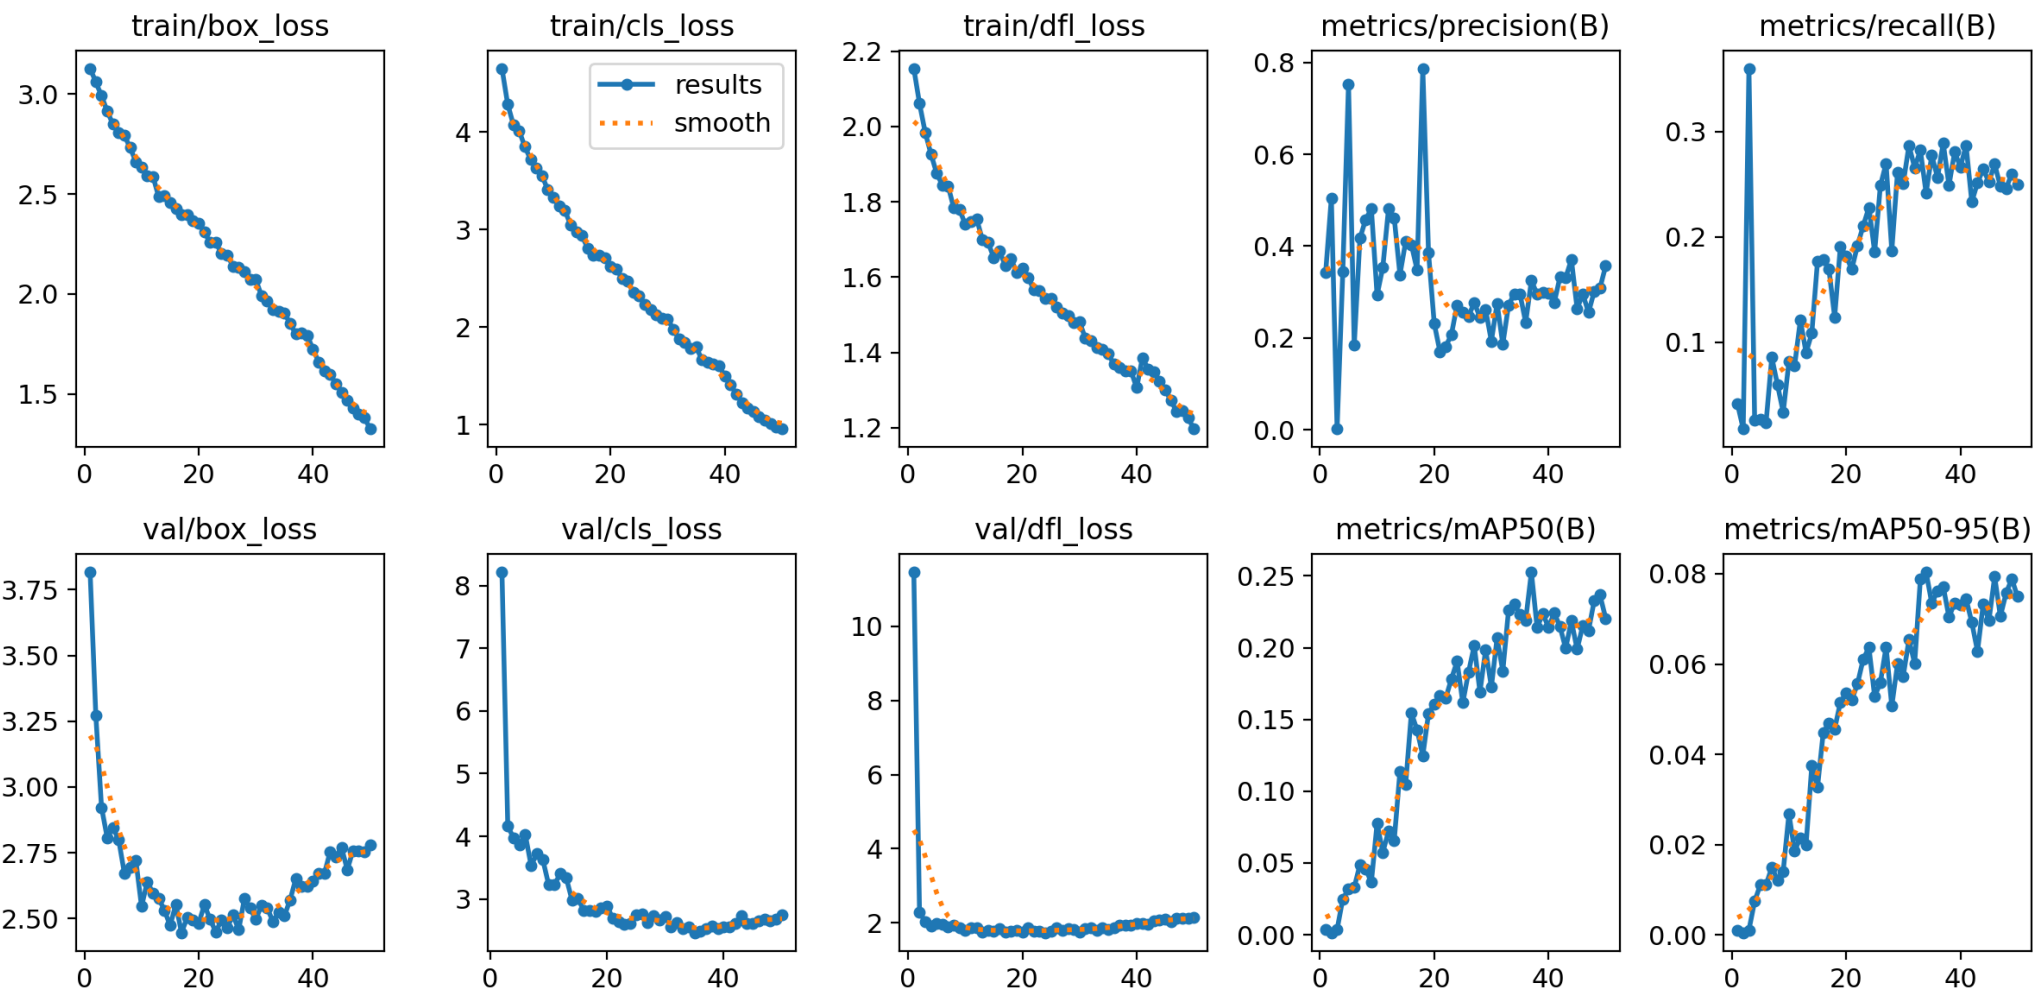
Model summary (fused): 268 layers, 68130309 parameters, 0 gradients, 257.4 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% 11
all	348	204	0.296	0.242	0.23	0.0802
elbow positive	348	29	0.0521	0.0345	0.0265	0.0094
fingers positive	348	48	0.389	0.146	0.204	0.0617
forearm fracture	348	43	0.538	0.465	0.438	0.175
humerus	348	36	0.442	0.583	0.593	0.204
shoulder fracture	348	20	0.111	0.15	0.0631	0.0137
wrist positive	348	28	0.244	0.0714	0.0568	0.0166


Speed: 0.1ms preprocess, 5.7ms inference, 0.0ms loss, 1.3ms postprocess per image




# 6.1 Results




# 6.1 Results



elbow positive  
fingers positive  
forearm fracture  
humerus  
shoulder fracture  
wrist positive  
all classes 0.26 at 0.155

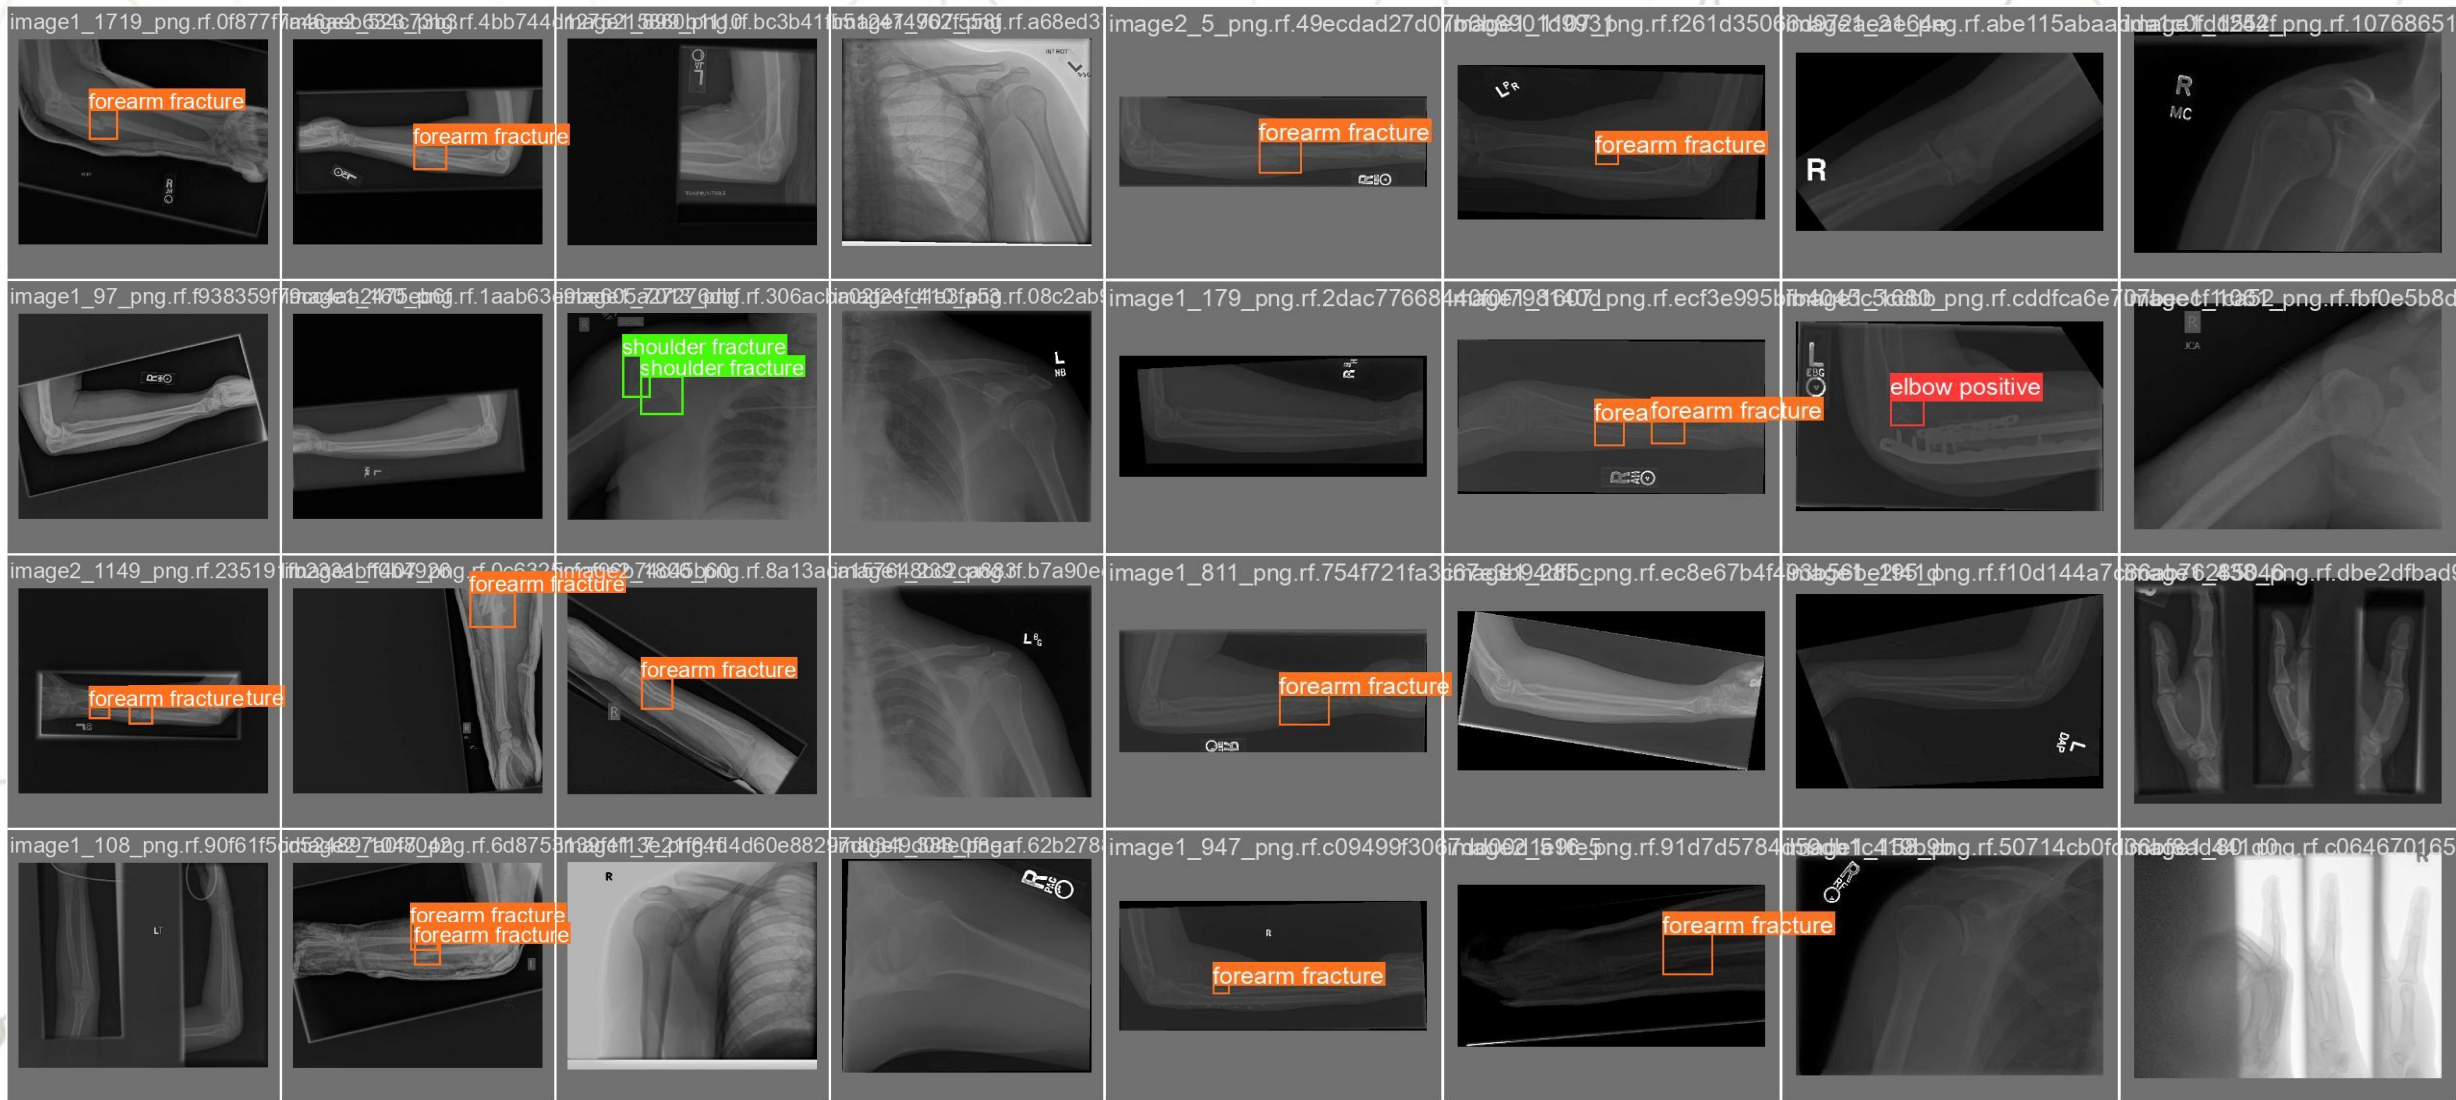


elbow positive 0.026  
fingers positive 0.204  
forearm fracture 0.438  
humerus 0.593  
shoulder fracture 0.063  
wrist positive 0.057  
all classes 0.230 mAP@0.5





# 6.1 Results





## 6.2 Results with Streamlit and FastAPI

### Fracture Bone Detection

Choose a file



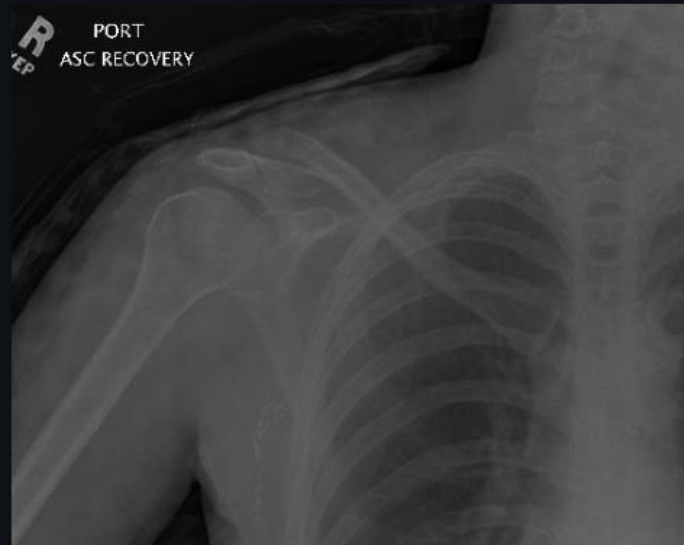
Drag and drop file here

Limit 200MB per file • JPG, JPEG, PNG

Browse files

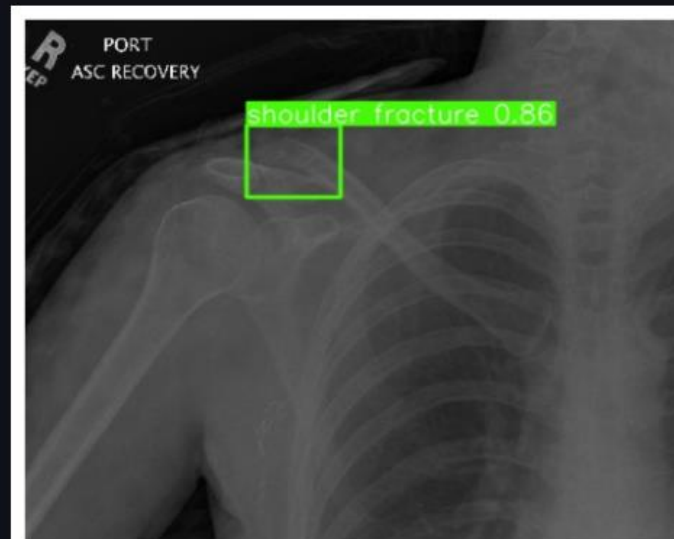


image1\_0\_png.rf.bc222874fa14e205b638e87ddad52b9d.jpg 15.3KB



Uploaded file

Prediction: YoloV8: Object Detection



# 7. YOLO-V8 WITH INSTANCE SEGMENTATION

The YoloV8-SEG model is a pretrained model downloaded from the library “Ultralytics”.

The YoloV8 was finetuned to bone fracture instance segmentation.

```
# Run inference on an image with YOLO
!yolo detect train data='./data.yaml' model=yolov8n-seg.pt epochs=90 imgsz=320 device=0
```

Downloading <https://github.com/ultralytics/assets/releases/download/v8.1.0/yolov8n-seg.pt> to 'yolov8n-seg.pt'...

100% 6.73M/6.73M [00:00<00:00, 24.7MB/s]

WARNING ⚠ conflicting 'task=detect' passed with 'task=segment' model. Ignoring 'task=detect' and updating to 'task=segment' to match model

Ultralytics YOLOv8.1.29 🚀 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4, 15102MiB)

**engine/trainer:** task=segment, mode=train, model=yolov8n-seg.pt, data=./data.yaml, epochs=90, time=None, patience=100, batch=16, imgsz=320, s

2024-03-20 22:10:23.819812: E external/local\_xla/xla/stream\_executor/cuda/cuda\_dnn.cc:9261] Unable to register cuDNN factory: Attempting to

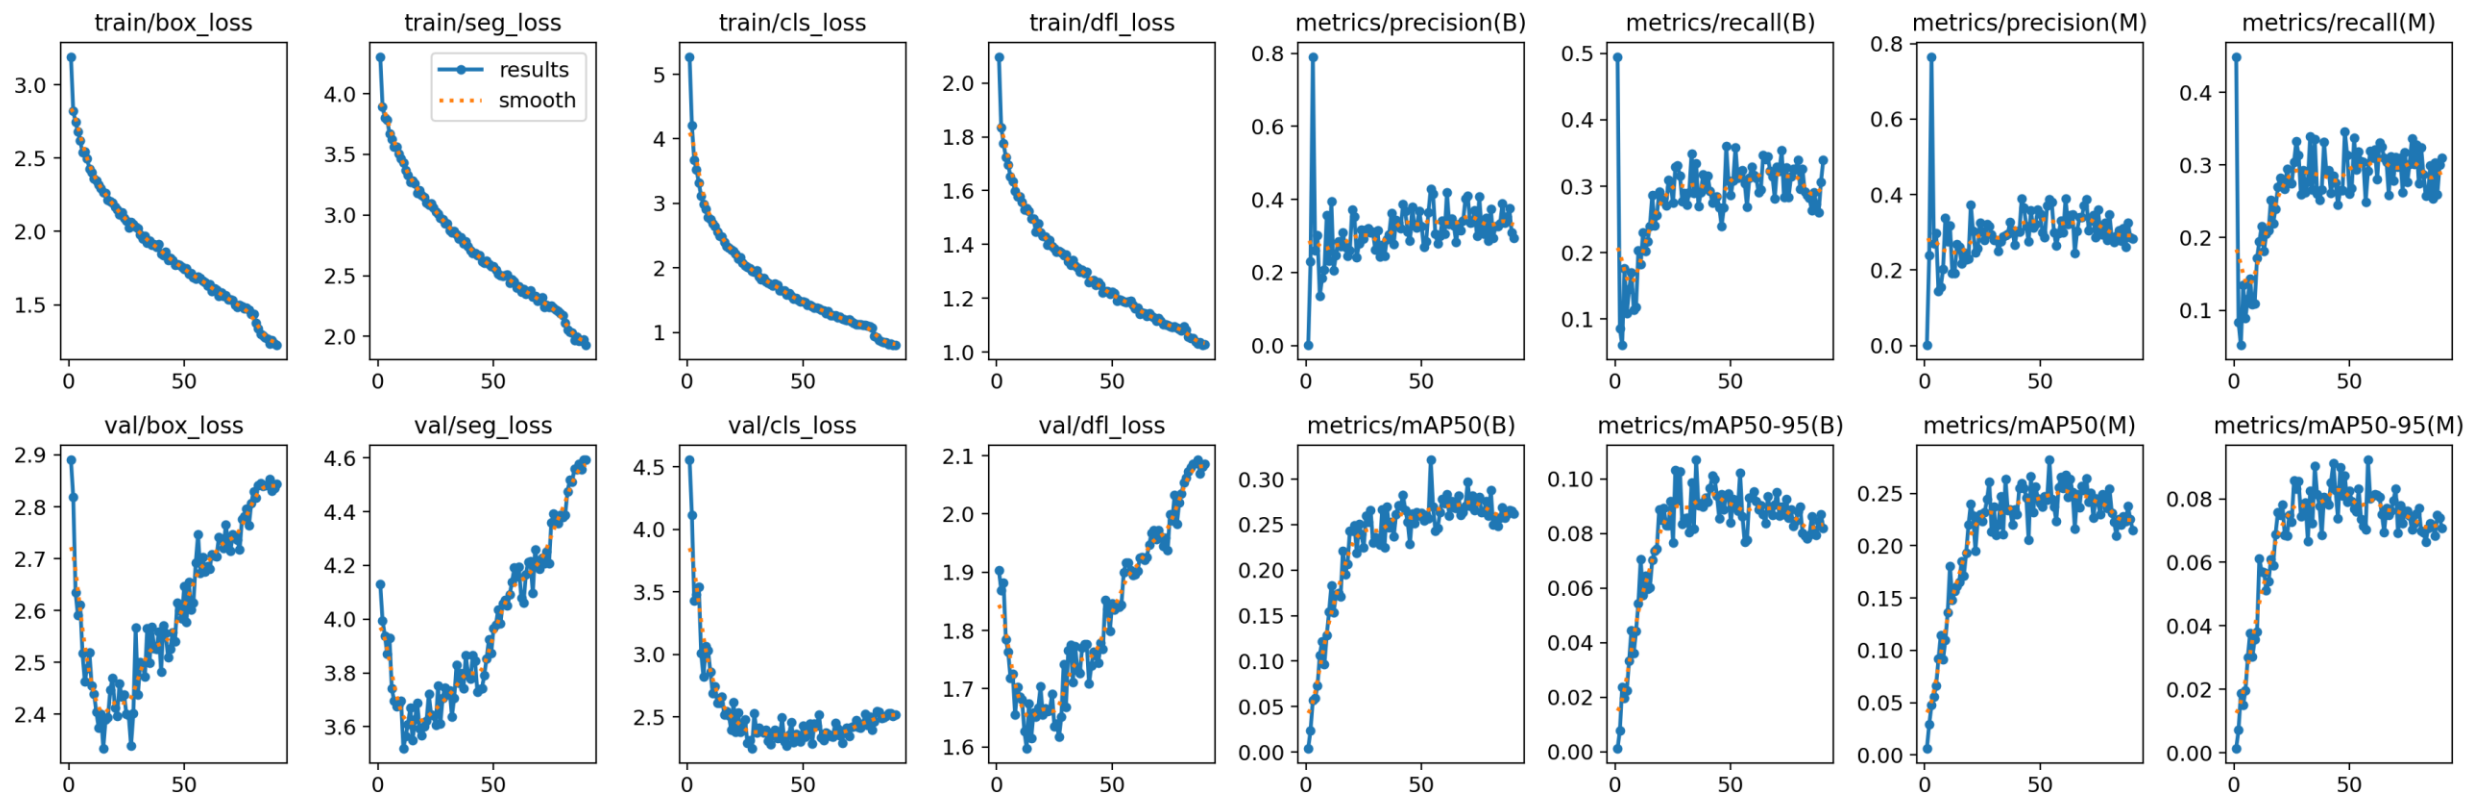
2024-03-20 22:10:23.819872: E external/local\_xla/xla/stream\_executor/cuda/cuda\_fft.cc:607] Unable to register cuFFT factory: Attempting to

2024-03-20 22:10:23.832882: E external/local\_xla/xla/stream\_executor/cuda/cuda\_blas.cc:1515] Unable to register cuBLAS factory: Attempting to



Overriding model.yaml nc=80 with nc=7

	from	n	params	module	arguments
0	-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1	-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2	-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4	-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8	-1	1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9	-1	1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10	-1	1	0	torch.nn.modules.unpooling.Unsample	[None, 2, 'nearest']

# 7.1 Results

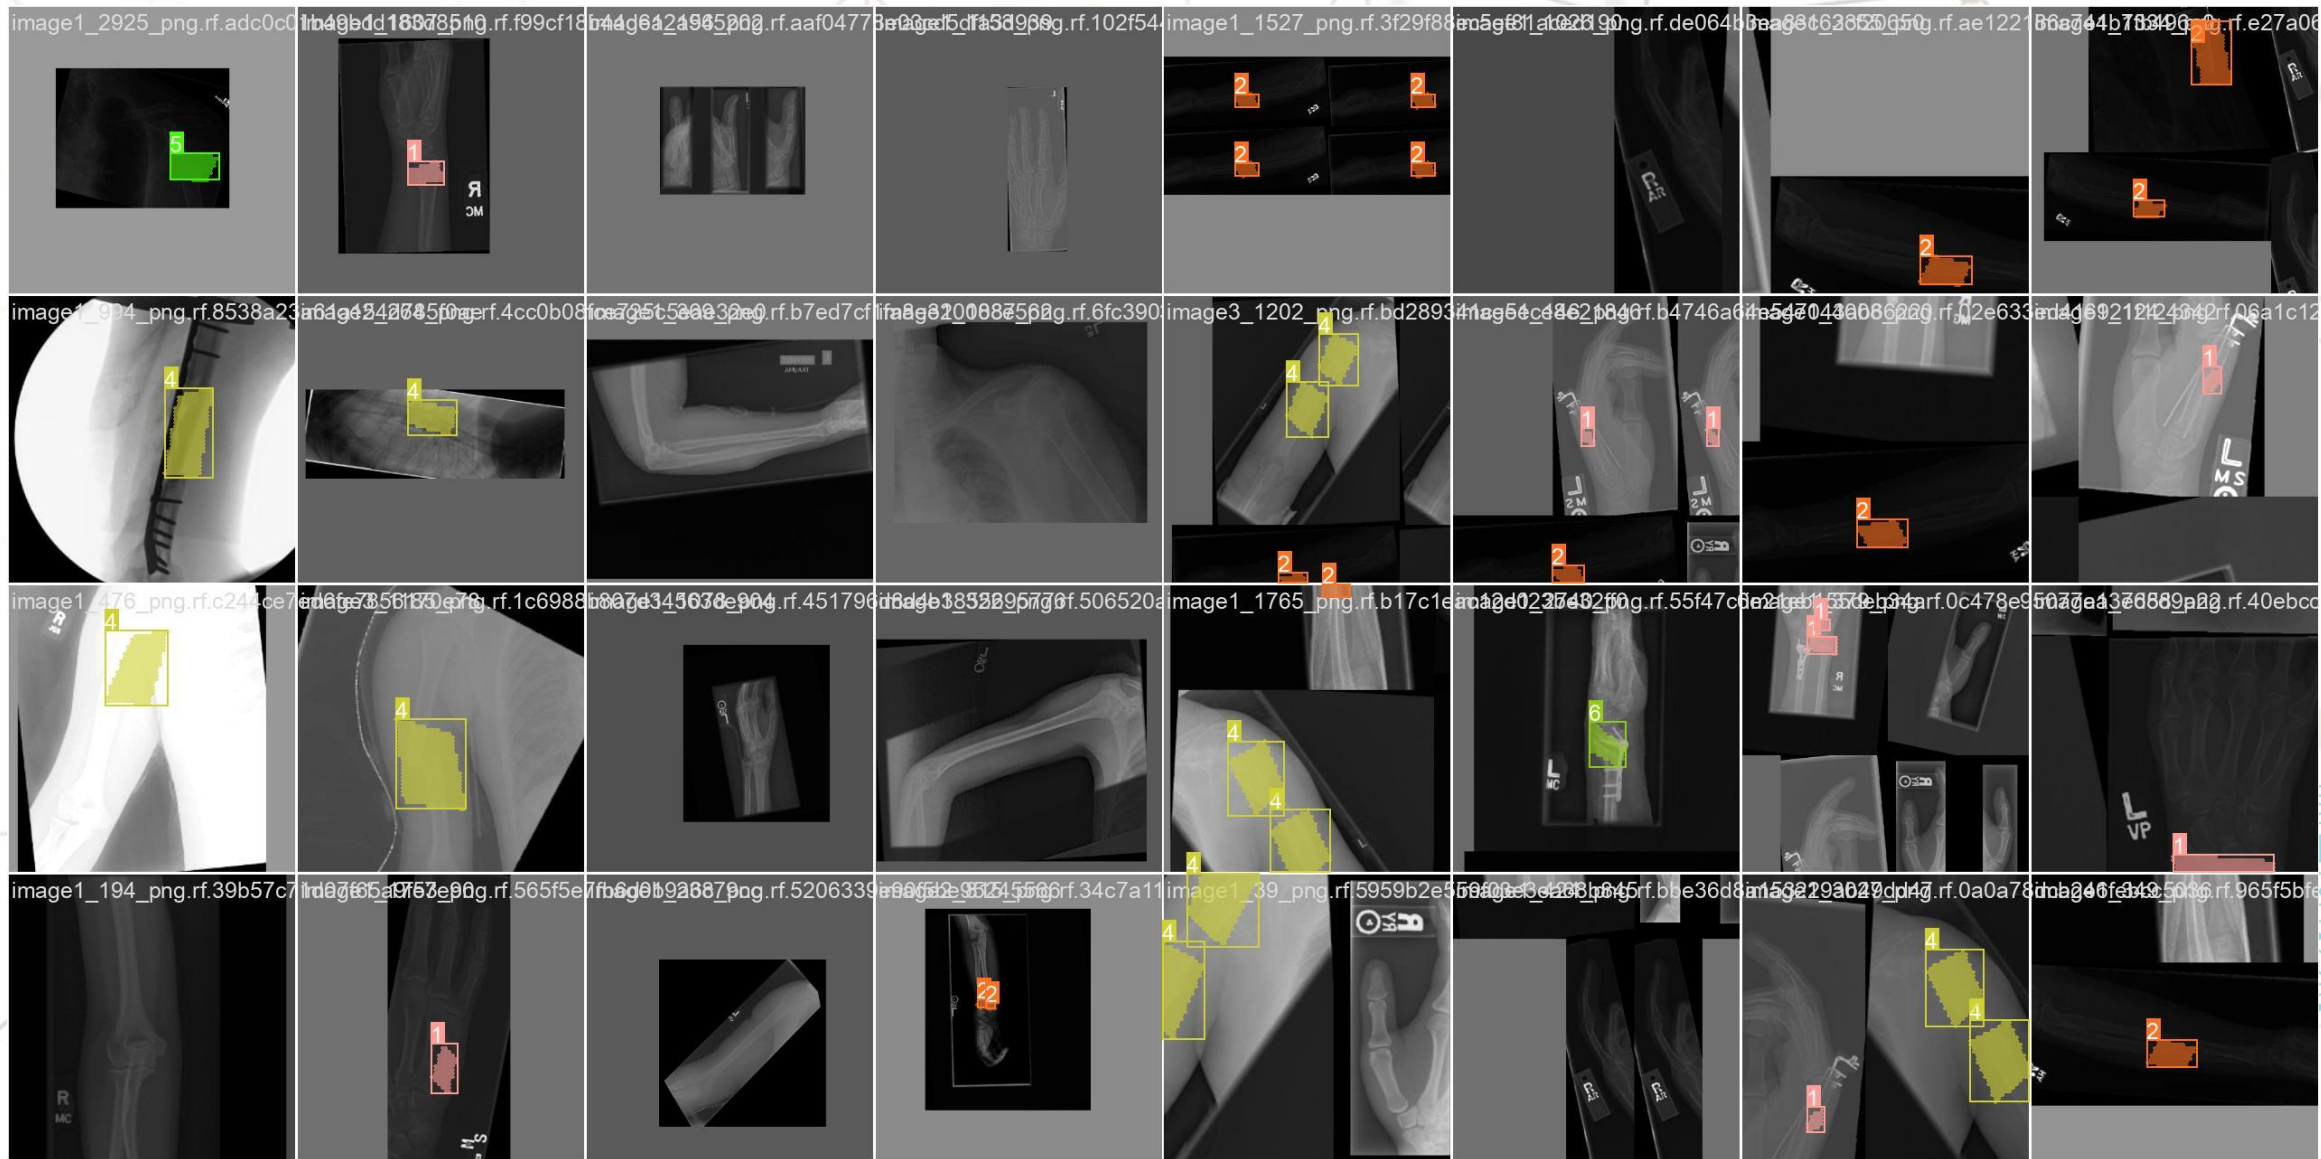


# 7.1 Results





# 7.1 Results





## 7.3 Results with Streamlit and FastAPI

### Fracture Bone Detection

Choose a file



Drag and drop file here

Limit 200MB per file • JPG, JPEG, PNG


Browse files



image1\_0\_png.rf.bc222874fa14e205b638e87ddad52b9d.jpg 15.3KB

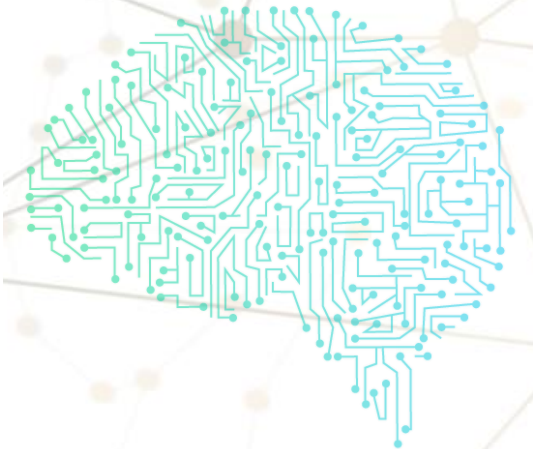
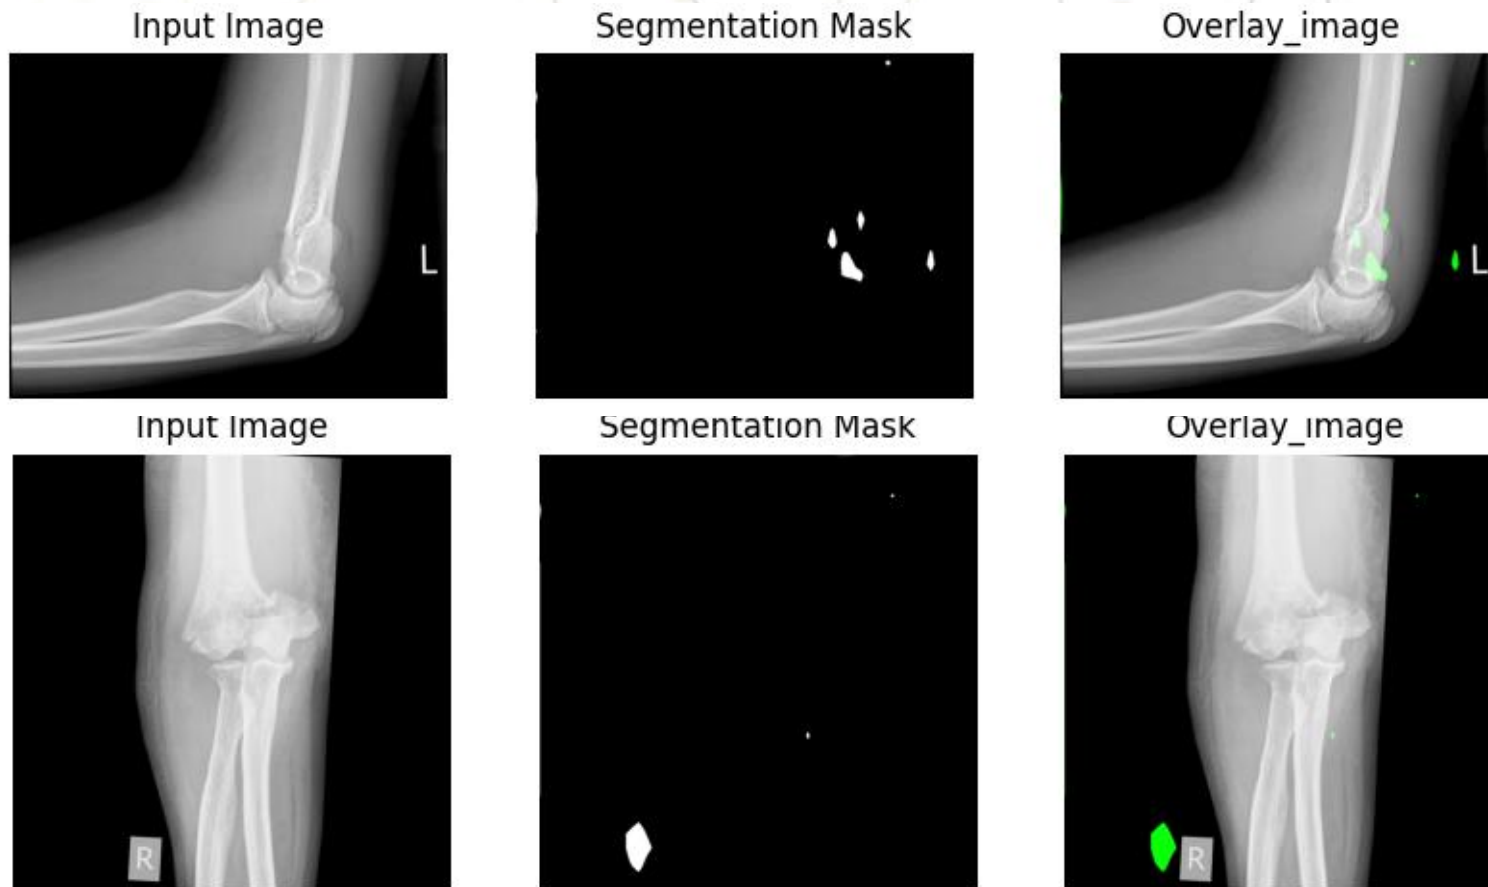


Prediction: YoloV8: Segmentation



# 8. Conclusions

1. Resnet struggle to make instance segmentation as is shown bellow:



# 8. Conclusions

2. Using Resnet with FastAPI has a lower response time in the server.
3. Using Yolo model has a relative higher response time in the server using streamlit and FastAPI.
4. Segmentation using YOLO could be improve by fine-tuning the model for more epochs or with more data for the training.
5. The deployment of the models on the server was successfully completed.

